

Greedy Algorithms

- ▶ Usually applied to **Optimization Problems**:
Find a configuration that minimizes or maximizes an objective function.
- ▶ A greedy algorithm...
 - ▶ Builds up a solution in small steps
 - ▶ At each step, makes its decision on some local criterion, myopically ignoring the global state
- ▶ Greedy algorithms that implement a greedy heuristic are easy. However,...
- ▶ Greedy algorithms that actually obtain a global optimum often do not exist.
- ▶ Greedy algorithms are usually simple, but the proof of correctness may not be so simple.
- ▶ We will give several examples of greedy algorithms.

Greedy Algorithm 1: Fractional knapsack problem

- ▶ Informal description:
 - ▶ There are n objects. Each object has a given weight and a given value.
 - ▶ There is a knapsack that can hold a given maximum weight.
 - ▶ Find the most valuable collection of objects that can fit into the knapsack.
- ▶ In the **fractional** version of the problem, we are allowed to break items into fractional portions.
- ▶ In the **zero-one** version of the problem, we either include all of each object or none of it.
- ▶ For now, we will consider just the fractional version.

Fractional knapsack problem: Formal problem description

- ▶ Inputs:
 - ▶ For each object ($i = 1, \dots, n$): **weight** = w_i , **value** = v_i
 - ▶ **capacity of knapsack** = W
- ▶ Output: For each object i :
 - ▶ x_i = the **weight of object i that we will put in the knapsack**
- ▶ The goal is to **maximize** the quantity

$$\sum_{i=1}^n \left(\frac{x_i}{w_i} \right) \cdot v_i$$

subject to the conditions

$$0 \leq x_i \leq w_i, i = 1, \dots, n$$

and

$$\sum_{i=1}^n x_i \leq W.$$

Fractional knapsack problem: Example

Suppose we have two objects:

- ▶ $w_1 = 80, v_1 = 80$
- ▶ $w_2 = 25, v_2 = 50$
- ▶ $W = 100$

Optimal solution:

- ▶ Take all of object 2 (weight = 25, value = 50)
- ▶ Take 75 weight units of object 1 (weight = 75, value = 75)
- ▶ Total weight = 100, total value = 125

Solution of Fractional Knapsack Problem

- ▶ Define the **value density** of object i to be v_i/w_i .
- ▶ Use object of highest value density in knapsack, then next highest, etc. Continue until knapsack is filled.
- ▶ Formal proof of correctness is in [GT] (page 288).

Implementation detail:

- ▶ Sort items in decreasing order of value density:
 - ▶ Algorithm Runs in $O(n \log n)$ time.
- ▶ Alternatively, implement using a max-heap. (Key of each object is its value density.)
 - ▶ Algorithm runs in $O(n + k \log n)$ time, where k is the number of objects fully or partially added to the knapsack.


Greedy Algorithm 2: Task scheduling problem

- ▶ Given n tasks
- ▶ Each task has a start time s_i and a finishing time f_i
- ▶ Each task must run on some machine. It will occupy that machine from its start time to its finishing time.
- ▶ Each machine can only work on one task at a time.
- ▶ Find the minimum number of required machines.

Task scheduling problem example

Example: 8 tasks:



$\{(12, 14), (5, 9), (0, 6), (6, 10), (1, 4), (3, 8), (8, 12), (2, 13)\}$



Greedy Solution to Task scheduling problem

- ▶ Process tasks according to start times, earliest start time first
- ▶ For each task T_i :
 - ▶ If we have a machine that can handle T_i , schedule T_i on that machine.
 - ▶ Otherwise, allocate a new machine, and schedule T_i on this new machine.

Solution to Task scheduling example



4 machines are sufficient

4 machines are necessary (so solution is optimal)

Greedy Algorithm 3: Single-machine scheduling


([GT] problem A-10.2)

- ▶ Given n tasks and only 1 machine.
- ▶ Each task has a start time s_i and a finish time f_i .
- ▶ If a task is selected, it will occupy the machine from its start time to its finishing time.
- ▶ The machine can only work on one task at a time.
- ▶ Find the maximum number of tasks that can be run.

Single-Machine scheduling problem example

Example: 8 tasks (same tasks as previous problem):

$\{(12, 14), (5, 9), (0, 6), (6, 10), (1, 4), (3, 8), (8, 12), (2, 13)\}$





Greedy Solution to Single-Machine scheduling problem

- ▶ Process tasks according to **finishing** times, earliest finishing time first.
- ▶ For each task T_i :
 - ▶ If task T_i does not conflict with tasks already chosen, select it.
 - ▶ Otherwise, do not select task T_i .
- ▶ [Kleinberg-Tardos] gives correctness proof and presents some plausible alternative greedy heuristics that do not work.
- ▶ This algorithm is optimal because the set of tasks it chooses "stays ahead" of any other set of tasks.

[Kleinberg-Tardos] Jon Kleinberg and Éva Tardos, Algorithm Design, Addison-Wesley, 2006.

Solution to Single-Machine scheduling Example



Choose 3 jobs: 5, 2, and 1

Final notes on Greedy Algorithms

- ▶ Greedy algorithms appear in various contexts. For example:
 - ▶ Graph problems: shortest paths, minimum spanning tree
 - ▶ String compression: Huffman coding
- ▶ Given an optimization problem, it is always worth asking: will a greedy algorithm work?
 - ▶ If the answer is yes, you have an easy solution.
 - ▶ If the answer is no (which it often is), seeing why the greedy approach fails may give you some insight into the structure of the problem.