



Graphs/Graph Algorithms


A **graph** is a set of objects (called **vertices**, **nodes**, or **points**) and a set of pairs of objects (called **edges** or **lines**)




Ring



Star





Wheel




Complete Graph

Directed vs. Undirected graphs

- ▶ Graphs can be **undirected** or **directed**.
- ▶ In a **directed graph**, or **digraph**, the edges have directions.



- ▶ It is possible to define a **mixed** graph with some edges undirected, others directed.
- ▶ For some purposes, we can replace an edge in an undirected graph by a pair of **antiparallel** directed edges



Examples of graphs (from [GT, Section 13.1])




- ▶ Visualizing binary relations.
 - ▶ Symmetric relations (e.g., “coauthored a paper with”)
 - ▶ Asymmetric relations (e.g., “inherits from”)
- ▶ City map
- ▶ Wiring/plumbing networks in buildings
- ▶ Flight network
- ▶ The Internet

Loops and Multiedges

- ▶ A **loop** (sometimes called a **self-loop**) is an edge with both endpoints the same



- ▶ Multiedges (also called **multiple edges**, or **parallel edges**)
 - ▶ A **multiedge** in an undirected graph is a pair of edges with the same endpoints.
 - ▶ A **multiedge** in a directed graph is a pair of edges with the same endpoints and the same direction.



Simple graphs

- ▶ A **simple graph** is a graph that has no loops and no multiedges.
- ▶ Unless we state otherwise, we will assume that all graphs are simple.


Notation

We write

- ▶ $G = (V, E)$ for the graph with vertex set V and edge collection E .
- ▶ n for the number of vertices of a graph
- ▶ m for the number of edges

Complete Graphs

- ▶ A **complete graph** is an undirected graph with an edge between every pair of vertices.
- ▶ A complete graph on n vertices has $m = \binom{n}{2}$ edges




Complete Graph on 7 vertices
($n = 7, m = 21$)

Some graph terminology

If $e = vw$ is an edge of a graph, we say:

- ▶ v (or w) is **incident** on e .
- ▶ e is **incident** on v (or on w).
- ▶ v (or w) an **endpoint** of e .
- ▶ v and w are **adjacent** (or **neighbors**)




- ▶ The **degree** of a vertex v is the number of edges incident on v .

Some additional terminology for digraphs

If $e = vw$ is an edge of a digraph, we say:

- ▶ v is the **origin**, or **tail**, of e .
- ▶ w is the **destination**, or **head**, of e .
- ▶ e is an **outgoing edge** of v .
- ▶ e is an **incoming edge** of w .



- ▶ The **indegree** of a vertex is the number of incoming edges.
- ▶ The **outdegree** of a vertex is the number of outgoing edges.

Some useful formulae:

1. In any graph

$$\sum_{v \in V(G)} \text{degree}(v) = 2m$$

2. In any directed graph

$$\sum_{v \in V(G)} \text{indegree}(v) = m$$

3. In any directed graph

$$\sum_{v \in V(G)} \text{outdegree}(v) = m$$

Paths


- ▶ A **path** from vertex v to vertex w is a sequence of edges

$$v_0 v_1, v_1 v_2, \dots, v_{k-1} v_k$$

1. $v_0 = v$
2. $v_k = w$

- ▶ The path is a **simple path** if

3. The vertices v_0, v_1, \dots, v_k are all distinct



Cycles

- ▶ In an **undirected** graph, a **cycle** is a set of edges

$$v_0 v_1, v_1 v_2, \dots, v_{k-2} v_{k-1}, v_{k-1} v_k$$

such that:

1. $v_0 = v_k$
 2. $k \geq 3$
- ▶ The cycle is **simple** if
 3. v_0, v_1, \dots, v_{k-1} are all distinct.
 - ▶ k is called the **length** of the cycle.
 - ▶ A graph is **acyclic** if it has no simple cycles.
 - ▶ In a **directed** graph, the definition is similar, but we can omit the requirement that $k \geq 3$, and we can omit the word “simple” in the definition of acyclic.


Cycles, continued

Usually when we write a cycle we just list the vertices, since the edges are implicit:

$$v_0, v_1, \dots, v_{k-1}, v_k$$

Sometimes we omit the last vertex:

$$v_0, v_1, \dots, v_{k-1}$$






So we can write: *AB,BC,CD,DA*

or *ABCD*

or just *ABCD*

Subgraphs


- ▶ A **subgraph** of $G = (V, E)$ is a graph $H = (V', E')$ such that $V' \subseteq V$ and $E' \subseteq E$




- ▶ H must be a valid graph. So every endpoint of an edge in E' must belong to V' .
- ▶ A **spanning subgraph** of G is a subgraph that contains all vertices of G .

Connected Graphs

- ▶ An undirected graph is **connected** if there is a path between any pair of vertices.




Connected



Not connected


Connected Components

- ▶ If a graph G is not connected, a **connected component** of G is a maximal connected subgraph of G .
- ▶ Connected components can also be defined in terms of an equivalence relation, reachability, on the vertices of G .
 - ▶ Vertex w is **reachable** from v if there is a path from v to w .
 - ▶ This relation is an equivalence relation (reflexive, symmetric, transitive)
 - ▶ The **connected components** are the equivalence classes of vertices, together with the edges connecting vertices in the equivalence class.




Trees


- ▶ A graph is **acyclic** if it has no simple cycles.
- ▶ A **tree** is a connected, acyclic graph.



A tree



Acyclic, but not
connected (a
forest)




Connected, but not
acyclic

Representations of graphs

1. Graph drawing
2. Edge List
3. Adjacency matrix
4. Adjacency list


1. Graph drawing



- ▶ Good for reasoning on paper, or for a GUI.

2. Edge List

- ▶ List of vertices, edges.
- ▶ Simple description for input, output.



$$V = (v_1, v_2, v_3, v_4)$$

$$E = (v_1 v_2, v_1 v_3, v_2 v_3, v_3 v_4)$$


3. Adjacency matrix

- ▶ Represent G with a matrix (2D array):

$$a_{i,j} = \begin{cases} 1 & \text{if } v_i v_j \text{ is an edge of } G \\ 0 & \text{otherwise} \end{cases}$$


- ▶ Space requirement = $\Theta(n^2)$
- ▶ Can modify to handle:
 - ▶ Directed graphs
 - ▶ Weighted graphs

	1	2	3	4
1	0	1	1	0
2	1	0	1	0
3	1	1	0	1
4	0	0	1	0



4. Adjacency list

- ▶ Vertices are stored in an array
- ▶ For each vertex, there is a pointer to a linked list describing its neighbors
- ▶ Space requirement = $\Theta(n + m)$
- ▶ Can modify to handle:
 - ▶ Directed graphs
 - ▶ Weighted graphs



Systematic traversal of graphs and digraphs

Two basic approaches:


- ▶ Depth-first search
- ▶ Breadth-first search

Examples of applications of depth-first-search:

- ▶ Testing whether a graph is connected
- ▶ Computing a spanning forest of a graph G (i.e., a subgraph that is a forest and contains every vertex of G)
- ▶ Computing the connected components of G
- ▶ Computing a path between two vertices v and w in a graph G (or reporting that no such path exists)
- ▶ Computing a cycle in an undirected graph G (or reporting that G is acyclic)



Depth-first search, Breadth-first search

- ▶ String/paint analogy
- ▶ **Depth-first search:** Follow path as far as possible, back up when dead end is reached.
e.g., *ABFGCDE HI KJ*
- ▶ **Breadth-first search:** Visit all neighbors of start vertex, then their neighbors, then neighbors of neighbors, etc.
e.g., *A BCDE FG JH KI*



Depth First Search (DFS) in an undirected graph

- ▶ Builds a DFS-forest (sometimes called a DFS-tree)
- ▶ Two kinds of edges:
 - ▶ **tree edge**, or **discovery edge**
 - ▶ **back edge**
- ▶ DFS-forest depends on choices of vertex ordering, ordering of neighbors about each vertex.
- ▶ Algorithms based on DFS return correct result for every result of these choices




Pseudocode for DFS in an undirected graph

- Initially, each edge and each vertex is unexplored

```
def DFS(G,v):  
    label v as explored  
    for all edges e incident on v:  
        if edge e is unexplored:  
            w ← opposite(v,e) // opposite endpoint of e from v  
            if vertex w is unexplored:  
                label e as a discovery edge  
                recursively call DFS(G,w)  
            else:  
                label e as a back edge
```

Analysis of the DFS algorithm

- ▶ $\text{DFS}(G, v)$ is called once per vertex
- ▶ Each edge is examined twice (once in each direction)
- ▶ Hence, total running time is $O(n + m)$.
- ▶ We may have to restart DFS multiple times to visit the entire graph.
- ▶ Running time will still be $O(n + m)$.




Connected-component labeling

Simple application of depth-first search in undirected graphs

Input: A graph G

Output: Each vertex $v \in V(G)$ is assigned a label, $L(v)$, such that two vertices are in the same connected component if and only if they have the same label.

Example:



Pseudocode for connected component labeling

- Initially, each vertex label is null, and each edge is unexplored.

Top level:

```

k = 0
for all vertices v in V(G):
    if v.label = null:
        k = k + 1
        DFSLabel(G,v,k)

```

Recursive Function to traverse and label the component:


```

def DFSLabel(G,v,k):
    v.label = k
    for all edges e incident on v:
        if edge e is unexplored:
            w ← opposite(v,e)
            mark e as explored
            if (w.label = null): DFSLabel(G,w,k)

```

Analysis: Runs in $O(m + n)$ time.

One application of connected-component labeling: Image processing



Binary image (black pixels are “on”, white pixels are “off”)

		1		2	2	2				
1	1		2	2	2	2				
	1		2	2	2	2				
			2	2	2	2				
			2	2	2	2				
			2		2	2				3
			2	2	2	2			3	3
			2		2	2			3	
			2	2	2	2			3	3
			2	2	2	2				3

Connected components of pixels in image that are “on”

Biconnected components, separation edges, separation vertices


(Material from [GT] Section 13.5. Read it!)

Let G be a connected graph.

- ▶ A **separation edge** is an edge whose removal causes G to become disconnected.
- ▶ A **separation vertex** is a vertex whose removal causes G to become disconnected.
- ▶ G is **biconnected** (or 2-connected) if for any two vertices $u, v \in V(G)$, there are at least two disjoint paths between u and v (i.e., two different paths that have no edges or vertices in common except for u and v).

Biconnected components



- Let G be a connected graph. A **biconnected component** (or **bicomponent**) of G is a subgraph G' such that either:
 - G' is biconnected, and adding any additional edges or vertices of G would force it to stop being biconnected; or
 - G' consists of a single separation edge and its two endpoints



Characterization of Biconnectivity

Lemma: Let G be a connected graph. The following are equivalent:

1. G is biconnected.
2. For any two vertices of G , there is a simple cycle containing them.
3. G does not have any separating vertices or separating edges.



Equivalence relations

Let $R(x, y)$ be a binary relation on a set of objects C . R is an equivalence relation if it satisfies the following three properties:

1. **Reflexive Property:** $R(x, x)$ is true for each x in C .
2. **Symmetric Property:** $R(x, y) = R(y, x)$ for each pair x and y in C .
3. **Transitive Property:** If $R(x, y)$ is true and $R(y, z)$ is true, then $R(x, z)$ is true, for every x , y , and z in C .

The **equivalence class** of an object x is the set of all objects y such that $R(x, y)$ is true.



Every element of C is in exactly one equivalence class.

Define a **link relation** on the edges of a graph G

Two edges e and f in $E(G)$ are **linked** if $e = f$ or if G has a simple cycle containing e and f .

Lemma: Let G be a connected graph. Then

1. The link relation forms an equivalence relation on the edges of G .
2. A bicomponent is the subgraph induced by the edges of an equivalence class of linked edges.
3. Edge e is a separation edge if and only if it is in a single-element equivalence class
4. Vertex v is a separation vertex if and only if it has incident edges in two different equivalence classes.




Preliminary biconnected components algorithm

Run DFS on G . Define an auxiliary graph F as follows:




- ▶ The vertices of F are the edges of G
- ▶ For every back edge e in G , let f_1, \dots, f_k be the discovery edges of G that form a cycle with e . F contains the edges $(e, f_1), \dots, (e, f_k)$.

Output the connected components of F . Each of these is a equivalence class of the link relation, and hence corresponds to a bicomponent of G .






Improved bicomponent algorithm

- ▶ The previous algorithm can require $\Omega(m \cdot n)$ time. We can improve this to $O(m + n)$.
- ▶ **Idea behind improved algorithm:** We don't need to compute F . We only need the connected components of F . So we build a spanning tree for each connected component of F (a spanning forest of F).



Improved bicomponent algorithm

1. Run DFS on G . Rank nodes of G according to order visited.
2. Add discovery edges of G to F . Mark discovery edges **unlinked**
3. Process back edges $e = (u, v)$ in order of rank of v . For each such e , let f_1, \dots, f_k be the discovery edges of G that form a cycle with e . Add edges $(e, f_k), (e, f_{k-1}), \dots, (e, f_j)$ to F , **stopping at first linked edge**. Mark edges $f_k, f_{k-1}, \dots, f_{j+1}$ **linked**



Pseudocode for Biconnected Component Algorithm

```




F ← an initially empty auxiliary graph
perform DFS traversal of G, starting at some vertex s
add each discovery edge f as a vertex in F, mark f "unlinked"
for each vertex v, in increasing rank order as visited in the
    DFS traversal
    for each back edge e=(u,v) with destination v
        add e as a vertex of the graph
        while u != v do
            Let f be the vertex in F corresponding to the
                discovery edge (parent(u),u)
            add the edge (e,f) to F
            if f is marked "unlinked" then
                mark f as "linked"
                u ← parent(u)
            else
                u ← v //exit while loop
compute the connected components of F

```

Note on the pseudocode

Note the sequencing:

1. Build DFS tree
2. Add discovery edges to F
3. Add back edges to F . Back edges are processed in the order in which their destination nodes were visited in DFS traversal.



Analysis of the Biconnected Component Algorithm

- ▶ DFS: $O(m + n)$ time.
- ▶ Building the auxiliary graph F : $O(m)$ time. Because...
 - ▶ Each iteration of the **while** loop causes an edge to be added to F .
 - ▶ When edge (e, f) is added to F :
 - ▶ Charge f if f is marked **unlinked**
 - ▶ Charge e if f otherwise
 - ▶ Each edge of G gets charged at most once
- ▶ Connected component analysis of F : $O(m)$ time.

So biconnected component analysis can be done in the same asymptotic time as connected component analysis.

Final note on Biconnected Component/Graph Vulnerability Analysis

As we have just seen, bicomponent analysis tells us:

- ▶ Whether G is 2-connected (i.e., whether any two vertices are joined by 2 disjoint paths).
- ▶ Whether G has a separating edge.
- ▶ Whether G has a separating vertex.



These questions can be generalized to any integer k :

- ▶ Are any two vertices joined by k mutually disjoint paths?
- ▶ Does G have a set of $k - 1$ edges whose removal makes G disconnected?
- ▶ Does G have a set of $k - 1$ vertices whose removal makes G disconnected?

These questions can be answered using **network flow** techniques.

Breadth First Search (BFS) in an undirected graph

- ▶ Start vertex is “level 0”
- ▶ “level $i + 1$ ” nodes are unexplored nodes that are neighbors of “level i ” nodes
- ▶ Process all nodes at one level before moving on to next level (special case: FIFO order)





Pseudocode for BFS in an undirected graph

```
def BFS(G,s):  
    Set all vertices and edges to unexplored  
    create an empty Queue Q[0]  
    insert s into Q[0]  
    i  $\leftarrow$  0  
    while Q[i] is not empty:  
        create an empty Queue Q[i+1]  
        for each vertex v in Q[i]:  
            for all edges e incident on v:  
                if edge e is unexplored:  
                    w  $\leftarrow$  opposite(v,e)  
                    if vertex w is unexplored:  
                        label e as a discovery edge  
                        insert w into Q[i+1]  
                    else:  
                        label e as a cross edge  
        i  $\leftarrow$  i+1
```

A useful property of BFS


- ▶ The level number of vertex v in BFS-tree rooted at s is the smallest number of edges in a path from s to v .



Analysis of the BFS algorithm

- ▶ Body of the outer **for** is performed n times
- ▶ Body of the inner **for** is performed $2m$ times
- ▶ Hence, total running time is $O(n + m)$.

As with DFS, we may have to restart BFS multiple times to visit the entire graph. Running time will still be $O(n + m)$.





Directed graphs

- ▶ Traversing a digraph (Brief mention of why DFS is more complicated in a directed graph than an undirected graph)
- ▶ Strong connectivity
- ▶ Directed acyclic graphs/topological orderings

DFS in a directed graph

- ▶ There are four kinds of edges (as opposed to only two in an undirected graph)
 - ▶ **tree edge**, or **discovery edge**
 - ▶ **back edge**: connects a vertex to an **ancestor** in the DFS tree
 - ▶ **forward edge**: connects a vertex to a **descendant** in the DFS tree
 - ▶ **cross edge**: connects a vertex to another vertex that is neither an ancestor nor a descendant in the DFS tree




Reachability


- ▶ Vertex v is **reachable** from vertex w in G if there is a path from w to v in G .
- ▶ In a digraph, the reachability relation is reflexive and transitive but not necessarily symmetric

Strong connectivity in digraphs

- ▶ A directed graph G is **strongly connected** if there is a path from every vertex in G to every other vertex in G .



Strongly Connected



Not Strongly Connected

Algorithm for testing strong connectivity in a digraph

1. Pick a start vertex s in G .
2. Run DFS in G , starting from s . If some vertex is not reachable from s , report NO and stop.
3. Let G^R be G with the direction of all edges reversed.
4. Run DFS in G^R , starting from s . If all vertices are reachable from s in G^R , report YES. Otherwise, report NO.

Analysis: Runs in $O(m + n)$ time.

Correctness:

- ▶ If either step 2 or step 4 says NO, G is not strongly connected.
- ▶ If both step 2 and step 4 say YES, G is strongly connected.

Proof: Pick any two vertices x and y in G . There is a path from x to y because ...



- ▶ Since step 4 said YES, there is a path from x to s
- ▶ Since step 2 said YES, there is a path from s to y
- ▶ Since reachability is transitive, there is a path from s to y

Directed Acyclic Digraphs (DAGs)

- ▶ A **Directed Acyclic Graph**, or **DAG**, is a directed graph with no cycles.
- ▶ **Examples:**
 - ▶ Courses in a degree program, with edges representing prerequisites.
 - ▶ Classes in C++ or Java, with edges representing inheritance
 - ▶ Tasks, with edges representing scheduling constraints

Topological ordering

- ▶ A **topological ordering** of a digraph G is a numbering of the vertices such that every edge is directed from a lower-numbered vertex to a higher-numbered vertex.



Topological ordering

Theorem: A directed graph has a topological ordering if and only if it is a DAG

Proof: see [GT], Section 13.4.4.

Proof Sketch:

- ▶ \Rightarrow : easy
- ▶ \Leftarrow : Assume G is a DAG.
 - ▶ G must have a vertex with indegree 0.
 - ▶ Let vertex v_1 be such a vertex.
 - ▶ Remove v_1 (and all incident edges) from G .
 - ▶ Result is a smaller DAG.




Topological Sorting

- ▶ The process of finding a topological ordering of a graph is called **topological sorting**
- ▶ Algorithm either
 - ▶ Finds a topological ordering, or
 - ▶ Reports that graph is not a DAG
- ▶ Algorithm repeatedly applies the following steps:
 1. Find a vertex v with indegree 0
 2. Assign v the next available label
 3. Delete v and all its outgoing edges
- ▶ Runs in $O(n + m)$ time using $O(n)$ additional space

Topological Sorting Pseudocode

```
def TopologicalSort(G):  
    L ← initially empty list of vertices  
    for each vertex v of G:  
        incounter(v) = indegree(v)  
        if incounter(v) = 0: add v to L  
  
    i ← 0  
    while L is not empty:  
        choose a vertex v in L and remove it from L.  
        i ← i+1  
        v.number ← i  
        for each edge e in v.outEdges:  
            w = opposite(v,e);  
            incounter(w) = incounter(w)-1;  
            if incounter(w) = 0: add w to L  
  
    if i == n: print the vertices and their numbers  
    else: print("G is not a DAG!!")
```


Topological Sorting Example



vertex	indegree
A	4 ³ 2100
B	100
C	21100
D	0
E	3 ² 2100
F	21100
H	0

$$L = [D, H]L = [H]L = [B]L =$$

Final Remark on Topological Sorting

Topological sorting is **non-deterministic**.

- ▶ At any step, there could be multiple vertices with indegree 0
- ▶ Any valid choice will lead to a valid topological ordering
- ▶ Note that this implies that a digraph may have many different topological orderings