

Weighted graphs

- ▶ A **weighted graph** is a graph that has a number $w(e)$ associated with each edge.
- ▶ Weights can be integers, rationals, reals, although in specific problems there may be restrictions.
- ▶ Weights can be positive, negative or 0, although in specific problems there may be restrictions.
- ▶ Weights can represent distance, cost, affinity, etc.

We will focus on two problems:

- ▶ **Shortest path problem**: Find path of minimum total weight, subject to specifications of the path
- ▶ **Minimum spanning tree problem**: Find spanning tree of minimum total weight

Shortest Paths


- ▶ Let G be a weighted graph. The **length** (or **weight**) of a path $P = e_0, \dots, e_{k-1}$ is the sum of the weights of the edges:

$$w(P) = \sum_{i=0}^{k-1} w(e_i)$$

- ▶ Let u, v be two vertices in V . A **shortest path** (or **minimum-length path** or **minimum-weight path**) from u to v is a path from u to v of minimum total weight.
- ▶ The **distance** from vertex u to vertex v , denoted $d(u, v)$, is the length of the shortest path from u to v if such a path exists.
- ▶ Note on Terminology
 - ▶ The weight of an edge is sometimes called its **cost** or its **length**
 - ▶ Similarly, the weight (or length) of a path is sometimes called its **cost**

Shortest paths: Notes on definitions

- ▶ If no path exists from u to v , we will say that $d(u, v) = +\infty$.
- ▶ Even if there is a path from u to v , there may not be a shortest path. (Because of **negative-weight cycles**, sometimes called **negative cycles**.)



Single Source Shortest Path Problem

- ▶ **Problem:** Given graph G and a vertex $v \in V(G)$, find the shortest path from v to every other vertex in $V(G)$.
- ▶ We will discuss two algorithms:
 1. Dijkstra's algorithm
 2. The Bellman-Ford Algorithm

Dijkstra's Algorithm

- ▶ G an undirected graph in which every edge weight is nonnegative
- ▶ “Weighted BFS” starting at v
- ▶ Grow a “cloud” (actually a tree) of points:
 - ▶ Initially, the cloud is empty
 - ▶ At each iteration:
 1. Choose the vertex outside the cloud that is closest to v .
 2. Add the chosen vertex to the cloud
- ▶ Example of a greedy algorithm

Overview of Dijkstra's algorithm


- ▶ Every vertex u has a **label**, $D[u]$.
- ▶ $D[u]$ stores the length of the best path from v to u that we have found **so far**
- ▶ Initially:
 - ▶ $D[v] = 0$ for our start vertex v
 - ▶ $D[u] = +\infty$ for all other vertices $v \neq u$
 - ▶ The vertex cloud C is empty (i.e., $C = \emptyset$)
- ▶ On each iteration:
 - ▶ Select a vertex u not in C with smallest $D[u]$ label
 - ▶ Put selected vertex u into C . (On first iteration, $u = v$).
 - ▶ Update $D[z]$ for each neighbor of u that is outside C (Because there may be a better path from v to z , via u , than we knew about before.)

```

if  $D[u] + w(u, z) < D[z]$  then
     $D[z] \leftarrow D[u] + w(u, z)$ 
  
```

This is called **edge relaxation**

Edge relaxation, shortest-path tree





```

if  $D[u] + w(u, z) < D[z]$  then
     $D[z] \leftarrow D[u] + w(u, z)$ 
     $z.parent = u;$ 

```

- ▶ Setting $z.parent$ organizes the points as a tree (the **shortest-path tree**), which allows us to find the shortest path after the algorithm completes.

Example of Dijkstra's Algorithm (Start vertex = A)



Pseudocode for Dijkstra's Algorithm

```
Algorithm DijkstraShortestPath( $G, v$ )
  for each vertex  $u$  in  $G$  such that  $u \neq v$  do
     $D[u] = +\infty$ 
   $D[v] = 0$ 
  for each vertex  $u$  in  $G$ 
     $u.parent = \text{null}$ 

  Put all vertices of  $G$  in a priority queue  $Q$ , using
    the labels  $D[]$  as keys
  while  $Q$  is not empty
     $u \leftarrow Q.removeMin()$  // Priority Queue Operation
    for each  $z$  adjacent to  $u$  such that  $z$  is in  $Q$  do
      if  $D[u] + w(u, z) < D[z]$  then
         $D[z] \leftarrow D[u] + w(u, z)$ 
         $Q.updateValue(z)$  to reflect new  $D[z]$  value
        // Priority Queue Operation
         $z.parent = u$ 

  return the collection of labels  $D[]$  and the parent values
```

Correctness of Dijkstra's algorithm

Follows from

- ▶ **Lemma:** When a vertex u is put in the “cloud” of known vertices (i.e., when u is removed from the priority queue), $D[u]$ represents the true distance from v to u

Correctness depends on there not being any negative edge weights.

Analysis of Dijkstra's Algorithm

Depends on implementation of priority queue

- ▶ Priority queue is abstract data type that supports removing the minimum, changing a value.
- ▶ Two possible implementations of a priority queue:
 1. **Heap**
 - ▶ Find/Remove minimum: $O(\log n)$
 - ▶ Update value: $O(\log n)$
 2. **Array**
 - ▶ Find/Remove minimum: $O(n)$
 - ▶ Update value: $O(1)$

Analysis of Dijkstra's Algorithm (continued)

- ▶ Operation counts:
 - ▶ `Q.removeMin()` performed n times
 - ▶ `Q.updateValue()` performed $O(m)$ times
- ▶ Standard implementation using heap:
 - ▶ Running time is $O((m + n) \log n)$
 - ▶ This simplifies to $O(m \log n)$ if G is connected
 - ▶ In terms of n only, this is $O(n^2 \log n)$ if G is simple
- ▶ Alternate implementation using an array:
 - ▶ Running time is $O((m + n^2))$
 - ▶ In terms of n only, this is $O(n^2)$ if G is simple
- ▶ Generally, heap implementation is better. However, if G is dense (in particular, if $m = \Omega(n / \log n)$), then the alternative implementation may be slightly better.
- ▶ **One-sentence summary:** Dijkstra's algorithm, using a heap, running on a simple connected graph, runs in $O(m \log n)$ time.

Finding the shortest path

After we have run Dijkstra's algorithm and computed the shortest-path tree, how do we find the actual shortest path from v to u ?

- ▶ We can read off the path in **reverse** order with the following loop:

```
while (u.parent  $\neq$  null)
    outputEdge(u,u.parent)
    u = u.parent
```


- ▶ We can read off the path in **forward** order by recursively calling `outputPath(u)`:

```
outputPath(u)
    if u.parent  $\neq$  null then
        outputPath(u.parent)
        outputEdge(u.parent,u)
```

Bellman-Ford Shortest-path Algorithm

- ▶ Directed graphs
- ▶ Allows negative edges (unlike Dijkstra)
- ▶ Does not allow negative cycles
- ▶ Based on “iterative relaxation” of edge weights
- ▶ Repeatedly cycle through edges

Bellman-Ford example



Pass 1: Pass 2: Pass 3: Pass 4: Pass 5: **Done**

$\Rightarrow DF$
 $\Rightarrow CD$
 $\Rightarrow AE$
 $\Rightarrow BC$
 $\Rightarrow AD$
 $\Rightarrow AB$
 $\Rightarrow EF$

u	$D[u]$
A	0
B	$+\infty$ 33
C	$+\infty$ 66
D	$+\infty$ 101099
E	$+\infty$ 22
F	$+\infty$ 87766

Pseudocode for Bellman Ford Algorithm

Algorithm BellmanFordShortestPath(G, v)

for each vertex u in G such that $u \neq v$ do

$D[u] = +\infty$

$D[v] = 0$

for $i \leftarrow 1$ to $n-1$ do

 for each edge $e=(u,z)$ in G do

 if $D[u] + w(u,z) < D[z]$ then

$D[z] \leftarrow D[u] + w(u,z)$

if $D[z] \leq D[u] + w(u,z)$ for all edges $e = (u,z)$ then

 return the collection of labels $D[]$

else

 return "G contains a negative-weight cycle"

Correctness of Bellman-Ford algorithm

- ▶ For any i : After i iterations, $D[u]$ is \leq the length of the shortest path from v to u that has at most i edges.
- ▶ So after $n - 1$ iterations, $D[u]$ is \leq the length of the shortest path from v to u that has at most $n - 1$ edges.
- ▶ Since no simple path can have more than $n - 1$ edges, the Bellman-Ford algorithm finds the shortest path provided there are no negative-weight cycles.

Analysis of Bellman-Ford algorithm

- ▶ $n - 1$ (i.e. $O(n)$) iterations of outer loop
- ▶ Each loop requires $O(m)$ relaxation tests
- ▶ Hence, total running time is $O(m \cdot n)$.

An optimization:


- ▶ We can stop after an iteration where no $D[u]$ values change.
(So we don't necessarily need to run $n - 1$ iterations.)

The Minimum Spanning Tree Problem

A **spanning tree** in a graph G is a free tree T such that



- ▶ T is a subgraph of G
- ▶ Every vertex of G is also a vertex of T .

A graph may have many different spanning trees



Minimum Spanning Tree



A **minimum spanning tree** (MST) in a weighted undirected graph is a spanning tree of minimum total edge weight.



Minimum Spanning Tree

If we assume that the weights represent costs:

A minimum spanning tree represents the cheapest way to connect all nodes of a graph, using edges of the graph.



Applications:


- ▶ Network layout
- ▶ Heuristics, subroutines for more complicated graph problems
 - ▶ Example: 2-Approximation for Traveling Salesman Problem with triangle inequality.

Uniqueness of Minimum Spanning Tree

- ▶ If all edge weights are distinct, the MST is unique.
- ▶ If two edge weights are the same, the MST may be unique, but it may not be.

A Crucial Fact about MST's

Let G be a weighted connected graph, and let V_1 and V_2 be a partition of $V(G)$ into two disjoint nonempty sets. Let e be an edge of minimum weight from among those edges with one endpoint in V_1 and the other endpoint in V_2 . Then there is a minimum spanning tree that has e as one of its edges.



Minimum Spanning Tree algorithms

We will discuss three algorithms for computing the MST:

- ▶ Prim-Jarník algorithm: grow MST as a tree from a “root vertex”
- ▶ Kruskal’s algorithm: Process edges in order of length
- ▶ Barůvka’s algorithm: Each connected component selects the smallest edge connecting it with another connected component


Prim-Jarník algorithm

- ▶ Grow a MST from a single cluster, starting from some “root vertex”
- ▶ Similar to Dijkstra’s algorithm (in fact, sometimes called the Prim-Dijkstra algorithm)
- ▶ Cloud of vertices C
- ▶ At each step:
 - ▶ Find smallest edge connecting a cloud vertex with a vertex not in the cloud.
 - ▶ Add this edge to the tree
 - ▶ Add endpoint outside cloud to the cloud.

Implementing the Prim-Jarník algorithm

- ▶ Each vertex u outside the cloud has a label $D[u]$ and another field, `u.parent`
- ▶ $D[u]$ stores the weight of the best edge connecting u to a cloud vertex.
- ▶ `u.parent` stores the other endpoint of this edge.
- ▶ When a vertex u is added to the cloud, the edge $(u, u.parent)$ is added to the tree.

Prim-Jarník Example (Root vertex = H)



Done!

Pseudocode for Prim-Jarník algorithm

Algorithm PrimJarnikMST(G)

pick an arbitrary root vertex v , set $D[v] = 0$

for each vertex u in G such that $u \neq v$ set $D[u] = +\infty$

for each vertex u in G set $u.parent = \text{null}$

Initialize $T \leftarrow \emptyset$

Put all vertices of G in a priority queue Q , using
the labels $D[]$ as keys

while Q is not empty

$u \leftarrow Q.removeMin()$ // Priority Queue Operation

 add vertex u to T

 if $u.parent \neq \text{null}$ add edge $(u, u.parent)$ to T

 for each z adjacent to u such that z is not in T do

 if $w(u, z) < D[z]$ then

$D[z] \leftarrow w(u, z)$

$z.parent = u$

$Q.updateValue(z)$ to reflect new $D[z]$ value

 // Priority Queue Operation

return the tree T

Analysis and Correctness of Prim-Jarník algorithm


- ▶ Correctness: Follows from Crucial Fact stated earlier
 - ▶ V_1 = vertices in the cloud
 - ▶ V_2 = vertices not in the cloud (i.e., in the priority queue Q)
- ▶ Analysis:
 - ▶ Runs in $O(m \log n)$ time, assuming we use a heap for the priority queue and the graph is connected

Kruskal's MST algorithm

Basic idea:

- ▶ Build MST in clusters
- ▶ Initially, each vertex is in its own cluster
- ▶ Process edges of graph in nondecreasing order of weight (from smallest to largest)
- ▶ Add an edge if its endpoints are in two different clusters
- ▶ Continue until $n - 1$ edges have been added

Example



	Edge	Wt.	Tree?
1	FG	2	YY
2	BG	3	YY
3	GH	4	YY
4	CD	5	YY
5	EF	5	YY
6	EH	6	NN
7	BC	6	YY
8	DE	7	NN

	Edge	Wt.	Tree?
9	AB	8	YY
10	AF	9	NN
11	JK	9	YY
12	CH	12	NN
13	IK	12	YY
14	IJ	14	NN
15	AI	19	YY

Done!

Implementation of Kruskal's algorithm

- ▶ Store edges in a priority queue, using the edge weight as the key
- ▶ We have to manage the clusters. We need to implement:
 - ▶ **Query:** Are vertex u and vertex v currently in the same cluster?
 - ▶ **Operation:** Merge two clusters

Pseudocode for Kruskal's algorithm

Algorithm KruskalMST(G)

for each vertex v in G

Define an elementary cluster $C(v)$, containing $\{v\}$ $//(*)$

Initialize a Priority Queue Q , storing the edges,
using the weights as keys

$T \leftarrow \emptyset$

while T has fewer than $n-1$ edges do

$(u,v) = Q.\text{removeMin}()$;

 Let $C(u) =$ the cluster containing u $//(*)$

 Let $C(v) =$ the cluster containing v $//(*)$

 If $C(u) \neq C(v)$ then $//(*)$

 Add edge (u,v) to T

 Merge $C(u)$ and $C(v)$ into a single cluster $//(*)$

- How do we implement “cluster management” functions at lines marked with $//(*)$?

Cluster Management

- ▶ Maintain each cluster as a linked list of vertices.
- ▶ Each vertex has an additional field `v.cluster`, indicating which cluster it belongs to
- ▶ When we merge two clusters, move elements of the **smaller** cluster into the **larger** cluster.

Analysis of Kruskal's Algorithm

- ▶ Priority queue operations: $O(m \log m)$, assuming heap implementation
 - ▶ Construct heap: $O(m)$
 - ▶ `removeMin()`: $O(m \log m)$ for all operations
 - ▶ $\leq m$ operations performed
 - ▶ Cost of each operation is $O(\log m)$ operation (which is actually $O(\log n)$ because G is simple),
- ▶ Cluster operations
 - ▶ Initializing clusters: $O(n)$
 - ▶ Merging clusters: $O(n \log n)$
 - ▶ Because each vertex gets moved between clusters at most $\log n$ times

So Kruskal's Algorithm runs in $O((m + n) \log n)$ time, which simplifies to $O(m \log n)$ time if G is connected.

Notes on Performance of Kruskal's Algorithm

- ▶ Limiting step is $O(m \log n)$ for heap operations. Cost of cluster operations match that step.
- ▶ If edges are already sorted, we can do better by using a faster but more complicated algorithm for cluster management, the **Union-Find** algorithm.
- ▶ We will come back to this later in these notes.


Correctness of Kruskal's Algorithm

- ▶ Follows from “Crucial Fact”
 - ▶ V_1 = cluster containing u
 - ▶ V_2 = all other vertices

Barůvka's MST algorithm

- ▶ Proceeds in “rounds.”
- ▶ Initially, each vertex is in its own cluster
- ▶ In each round:
 - ▶ Each cluster C selects the smallest edge in with one endpoint in C , the other endpoint outside C
 - ▶ All selected edges are added to MST.
 - ▶ Clusters are merged
- ▶ Continue until only one cluster remains

Baruvka's algorithm example



Round 1:

Cluster


{A}	ABAB
{B}	BGBG
{C}	CD CD
{D}	CD CD
{E}	EF EF
{F}	FG FG
{G}	FG FG
{H}	GH GH
{I}	IK IK
{J}	KJ KJ
{K}	KJ KJ

Selected Edge Cluster

{A, B, E, F, G, H}	BCBC
{C, D}	BCBC
{I, J, K}	AI AI

Round 2:

DONE!



Pseudocode for Barůvka's algorithm

Algorithm BarůvkaMST

```
Let T be a subgraph initially containing
    just the vertices of G (and no edges)
while T has fewer than  $n-1$  edges do
    for each connected component C of T
        find the smallest weight edge  $e=(u,v)$ 
            with  $u$  in C and  $v$  not in C
        add  $e$  to T (unless  $e$  is already in T)
return T
```


Implementation of Barůvka's algorithm

- ▶ Store T as an adjacency list.
- ▶ In each round:
 1. Label connected components (each vertex gets a label indicating to which component it belongs)
 2. Scan edges to find edge that will be selected by each component.

This takes $O(m)$ time per round

Analysis of Barůvka's algorithm

- ▶ Performance:
 - ▶ Since the number of components decreases by a factor of at least 2 each round, the number of rounds is $\leq \log n$,
 - ▶ The computation in each round is performed in $O(m)$ time
 - ▶ So running time of Barůvka's algorithm is $O(m \log n)$.
- ▶ Correctness:
 - ▶ follows from “Crucial Fact”, with $V_1 = C$, $V_2 = \text{everything else}$.
 - ▶ **Technical issue:** avoiding cycles. Number edges, use edge numbers to break ties.

The Union-Find Algorithm

Kruskal's algorithm

- ▶ Builds MST in clusters
- ▶ As we process each edge, we need to
 - ▶ Determine whether the two endpoints are in the same cluster
 - ▶ If not, add the edge to the MST and merge the two clusters
- ▶ Continue until only one cluster remains

We describe a very efficient algorithm for managing these clusters.

Set merging problem

The cluster management problem that arises in Kruskal's algorithm is a special case of the **dynamic equivalence problem**:

- ▶ Objects
- ▶ Sets of objects, with membership changing over time
- ▶ Queries: are two objects in the same set?

The special case that we need for cluster management in Kruskal's algorithm is called the **set merging problem**. The sets change as follows:

- ▶ Initially, each object is in its own set
- ▶ Over time, sets get merged. Once two sets get merged, they never get separated.

The set merging problem arises in other contexts. For example

- ▶ We are given some of the equivalence pairs in an equivalence relation, and we want to determine the equivalence classes.
- ▶ We are given the edges of a graph in some arbitrary order, and we want to construct the connected components without building the entire graph.

Example:

Suppose we have three sets:

set 1: $\{A, B, E\}$

set 2: $\{D, G\}$

set 3: $\{C\}$

A and C are not in the same set. But if we merged set 1 and set 3 and asked the question again, they would be in the same set.

The Union-Find algorithm

To solve the set merging problem, we will implement two primitive operations, called `union` and `find`.

1. `find(x)`: returns a temporary name for the set to which the object `x` currently belongs.

More precisely, `find(x)` returns a special element in the set to which `x` belongs. If `x` and `y` are in the same set, then `find(x)` and `find(y)` will return the same value.

2. `union(s,t)`: If `s` and `t` are the temporary names of two different sets, this command causes the two sets to be merged.

The Union-Find algorithm

Once we have the operations `union` and `find`:

- ▶ To determine whether objects x and y are in the same set:
`if find(x) == find(y) ...`
- ▶ To determine whether x and y are in the same set and, if not, to merge them:

```
s = find(x);  
t = find(y);  
if s  $\neq$  t  
    union(s,t);
```

Data representation

- ▶ Each set is a **tree** of objects (not necessarily a binary tree). Each object has a parent pointer, but no child pointer.
- ▶ A **find(x)** command returns the object at the root of the tree to which **x** currently belongs.
- ▶ A **union(x,y)** command is only valid if **x** and **y** are both roots of trees. It combines the two trees by either
 - ▶ making **x** the parent of **y**, or
 - ▶ making **y** the parent of **x**

We will discuss shortly how to make the choice.


Example




`union(A,B)`




`union(C,D)`



`union(C,E)`



`union(B,C)`



Union-Find


The algorithm as just stated is not very efficient.

For example, suppose there are $n + 1$ objects x_0, x_1, \dots, x_n . Suppose we first issue the following `union()` commands

```
union( $x_0, x_1$ )  
union( $x_1, x_2$ )  
...  
union( $x_{n-1}, x_n$ )
```

and then call `find(x_0)` repeatedly.

One possible result of the `union()` operations is as shown. Each call to `find()` would then require $\Theta(n)$ work.



Two simple improvements to Union-Find


We can make the Union-Find algorithm extremely efficient with two simple improvements:

1. **Weight balancing** on **union**.
2. **Path compression** on **find**.

Improvement 1: Weight balancing on union

When two trees are merged, the root of the **larger** tree (the tree with more nodes) becomes the root of the new merged tree.

Example:



Implementation note



Implementation of weight-balancing:

- ▶ Each node has a `weight` field.
- ▶ For a root node, the `weight` field contains the number of nodes in the tree rooted at that node.
 - ▶ Initially, each `weight` field is 1.
 - ▶ When two trees are merged, the `weight` field of the non-surviving root node is added to the `weight` field of the surviving root node.

Improvement 2: Path compression on find

On a **find(x)** operation, make every non-root node encountered on the path from **x** to the root a child of the root.

Example: **find(E)**:



With these two simple improvements, the union-find algorithm runs in “almost constant time per operation”

Analysis of the Union-Find Algorithm

- ▶ To describe the running time of the Union-Find algorithm, we first need to introduce a new function, called $\alpha(n)$ (“alpha”), that approaches ∞ as n gets large but does so extremely slowly.
- ▶ The function $\alpha(n)$ is the inverse of another function, $A(n)$, which approaches ∞ very quickly. $A(n)$ is defined as follows:

$$\begin{cases} A(0) &= 1 \\ A(n) &= 2^{A(n-1)} \end{cases} \quad \text{if } n > 0$$

$$A(0) = 1$$

$$A(1) = 2$$

$$A(2) = 4$$

$$A(3) = 16$$

$$A(4) = 65536$$

$$A(5) = 2^{65536} \approx 2 \times 10^{19728}$$

$$A(6) = 2^{2^{65536}}$$

The function $\alpha(n)$

$\alpha(n)$ = the smallest i such that $A(i) \geq n$

$$\alpha(1) = 0$$

$$\alpha(2) = 1$$

$$\alpha(n) = 2 \quad (2 < n \leq 4)$$

$$\alpha(n) = 3 \quad (4 < n \leq 16)$$

$$\alpha(n) = 4 \quad (16 < n \leq 65536)$$

$$\alpha(n) = 5 \quad (65536 < n \leq 2^{65536})$$

$$\alpha(n) = 6 \quad (2^{65536} < n \leq 2^{2^{65536}})$$

$\alpha(n)$ grows very slowly, but

$$\lim_{n \rightarrow \infty} \alpha(n) = \infty$$

Performance of Union-Find

It can be shown that

1. Any sequence of k union and find operations on a set of n objects, each initially in its own set, has a total cost of

$$O(k\alpha(n))$$

2. There are sequences of k union and find operations on a set of n objects, each initially in its own set, that do in fact require

$$\Omega(k\alpha(n))$$

So union-find has a cost per operation that is asymptotically worse than $\Theta(1)$, but is “almost” constant.

Note: This material can be found in [GT 4.2] and [CLRS 21]. [GT] proves a weaker analysis result. A proof of the analysis result stated here can be found in [CLRS].